

# Ornstein-Uhlenbeck process

June 14, 2016

## 1 Numerical *simulation* of the OU stochastic process

Let's consider the following stochastic differential equation that defines a Ornstein-Uhlenbeck (OU) process  $x(t)$ , whose (steady-state) mean and variance are  $\mu$  and  $\sigma^2$ , respectively:

$$\tau \frac{dx}{dt} = \mu - x + \sigma \sqrt{2\tau} \xi(t)$$

Abusing the notation of ordinary differential equations,  $\xi(t)$  is a *Gaussian white-noise* characterised by zero mean, unitary variance, and Dirac's Delta autocorrelation, with Gaussian probability density distribution as follows:

$$\langle \xi(t) \rangle = 0 \quad \langle \xi(t) \xi(t+T) \rangle = \delta(T) \quad p_\xi(Z) = (\sqrt{2\pi})^{-1} e^{-Z^2/2}.$$

$x(t)$  is then a Gauss-distributed non-stationary process, fully characterised by its time-varying mean  $m_x(t)$ , variance  $s(t)$  and autocovariance. The probability density distribution is then  $p_x(X) = (\sqrt{2\pi s(t)})^{-1} e^{-(X-m(t))^2/(2s(t))}$ , with

$$m(t) = \langle x(t) \rangle = x_0 e^{-(t-t_0)/\tau} + \mu \left(1 - e^{-(t-t_0)/\tau}\right) \longrightarrow \mu$$

$$s(t)^2 = \langle [x(t) - m(t)]^2 \rangle = \sigma^2 \left(1 - e^{-2(t-t_0)/\tau}\right) \longrightarrow \sigma^2$$

where  $x_0$  indicates the (deterministic) initial condition of  $x(t)$  at time  $t_0$ , and the autocovariance

$$Cov_x(t, T) = \langle [x(t) - m(t)] [x(t+T) - m(t+T)] \rangle = \sigma^2 e^{-|T|/\tau} (1 - e^{-2(t-t_0)/\tau}) \longrightarrow \sigma^2 e^{-|T|/\tau}$$

Numerically simulating the continuous-time process  $x(t)$  requires iterating the following expression

$$y(t + \Delta t) \approx (1 - \Delta t/\tau) y(t) + \mu \Delta t/\tau + \sigma \sqrt{2\Delta t/\tau} \hat{\xi}$$

where the (integration) time step  $\Delta t$  must be *small enough*, and where at every iteration  $\hat{\xi}$  is generated as a new realisation of a pseudo-random zero mean, unitary variance, uncorrelated Gauss-distributed number:

$$\langle \hat{\xi} \rangle = 0 \quad \langle \hat{\xi}_k \hat{\xi}_m \rangle = \delta_{k\ m} \quad p_{\hat{\xi}}(Z) = (\sqrt{2\pi})^{-1} e^{-Z^2/2}$$

with  $\delta_{k\ m}$  being the Kronecker's delta.

However, an alternative exact update rule is also available (Gillespie, 1992) that is valid for any finite value of  $\Delta t$ :

$$y(t + \Delta t) = e^{-\Delta t/\tau} y(t) + \mu \left(1 - e^{-\Delta t/\tau}\right) + \sigma \sqrt{1 - e^{-2\Delta t/\tau}} \hat{\xi}$$

It is easy to prove that, in the limit of small  $\Delta t$ , such an exact update rule reduces to the approximate expression above.

In Julia, one can easily generate a realisation of this OU process, containing  $N$  points, by three different style of code:

In [89]: # Let's first define the numerical parameters of our numerical example

```

μ = 50.;          # Steady-state mean of x(t)
σ = 6.;          # Steady-state standard deviation of x(t)
τ = 2.;          # Relaxation time [ms]

```

```

x0 = 0.;          # (Deterministic) initial condition

```

```

Δt = 0.02;       # Discretization time interval [ms]

```

```

N = 10000000;    # No. of samples to generate (beware of transients if x0 is not zero)

```

In [90]: # Let's preallocate the memory for three vectors, which will be used for t

```

x1 = zeros{Float64, N}; # Let's preallocate the memory for x1 - note
x2 = zeros{Float64, N}; # Let's preallocate the memory for x2 - note
x3 = zeros{Float64, N}; # Let's preallocate the memory for x3 - note

```

The first method does not exploit the true potential of Julia for numerical computation, as no *function* is defined.

In [91]: # The first method (slowest) is based on common-sense scripting

```

srand(359);      # Set a (fixed) random number seed - for reproducibility purposes

```

```

tmp1 = exp(-Δt/τ);          # Useful constant to slightly reduce the number of exp calls
tmp2 = σ * sqrt(1-exp(-2*Δt/τ)); # Useful constant to slightly reduce the number of sqrt calls
tmp3 = μ * (1 - exp(-Δt/τ)); # Useful constant to slightly reduce the number of exp calls

```

```

x1[1] = x0;                # We take care of the initial condition

```

```

@time for i=2:N,
    x1[i] = x1[i-1] * tmp1 + tmp3 + tmp2 * randn(); # exact iterative rule
end

```

```

6.015551 seconds (110.00 M allocations: 1.937 GB, 6.32% gc time)

```

The second method defines a function, calls it with *minimal size* to have it compiled, and then fires it up with full-sized problem.

```

In [92]: # The second method (~30-50 times faster) is based on Julia's just-in-time compilation

# We first define a function, which changes one of its input argument (for in-place)
function simulate_ou!(x, N, x0, Δt, μ, σ, τ)
    tmp1 = exp(-Δt/τ); # Useful to slightly reduce the n. of allocations
    tmp2 = σ * sqrt(1-exp(-2*Δt/τ)); # Useful to slightly reduce the n. of allocations
    tmp3 = μ * (1 - exp(-Δt/τ)); # Useful to slightly reduce the n. of allocations

    x[1] = x0; # We take care of the initial condition
    for i=2:N,
        x[i] = x[i-1] * tmp1 + tmp3 + tmp2 * randn();
    end
end

# We call the function once, with N = 2, so that it is compiled once for a real call
simulate_ou!(x2, 2, x0, Δt, μ, σ, τ);

# We then run the real call to generate N samples
srand(359); # Set a (fixed) random number seed - for reproducibility purposes

@time simulate_ou!(x2, N, x0, Δt, μ, σ, τ);

0.127130 seconds (4 allocations: 160 bytes)

```

The third method is the vectorised version of the first and employs the function `filt()`. It consumes more memory and it is slower.

```

In [93]: # The third method (slightly slower) is based on Julia's own DSP and filter functions

srand(359); # Set a (fixed) random number seed - for reproducibility purposes

tmp1 = exp(-Δt/τ); # Useful to slightly reduce the n. of allocations
tmp2 = σ * sqrt(1. - exp(-2*Δt/τ)); # Useful to slightly reduce the n. of allocations
tmp3 = μ * (1. - exp(-Δt/τ)); # Useful to slightly reduce the n. of allocations

b = [0, 1]; # Refer to the documentation of "filt!()" - these are the filter coefficients
a = [1, -tmp1]; # to efficiently compute a discrete-time iterative filter

@time filt!(x3, b, a, tmp3 + tmp2 * randn(N), [x0]);

0.310261 seconds (13 allocations: 228.882 MB, 10.60% gc time)

```

By virtue of the fixed random number seed initialisation, produces exactly the same sequence of numerical values, as confirmed below (apart from very minor rounding errors).

```

In [94]: # Let's see whether the numerical sequences are identical (result should be true)

```

```

test1 = 1./N * sqrt(sum(x1.*x1 - x2.*x2));
test2 = 1./N * sqrt(sum(x1.*x1 - x3.*x3));
test3 = 1./N * sqrt(sum(x2.*x2 - x3.*x3));

println("The mean square error comparing x1 and x2 is $test1");
println("The mean square error comparing x1 and x3 is $test2");
println("The mean square error comparing x2 and x3 is $test3");

using PyPlot
fig = figure("Sample realisation",figsize=(6,5)); # Create a new figure, w

index = (N-300):N;
plot(index, x1[index], "r", label="Method 1");
plot(index, x2[index], "g", label="Method 2");
plot(index, x3[index], "b", label="Method 3");

ax = gca() # Get a "handle" on the current axes
ax[:set_ylim](( $\mu-2*\sigma$ ), ( $\mu+2*\sigma$ ));

xlabel("index - last 300 points"); # Label of the horizontal axis
ylabel("x"); # Label of the vertical axis
title("OU - sample realisation (last 300 points)"); # Label of the title,
grid("on"); # "Grid" on

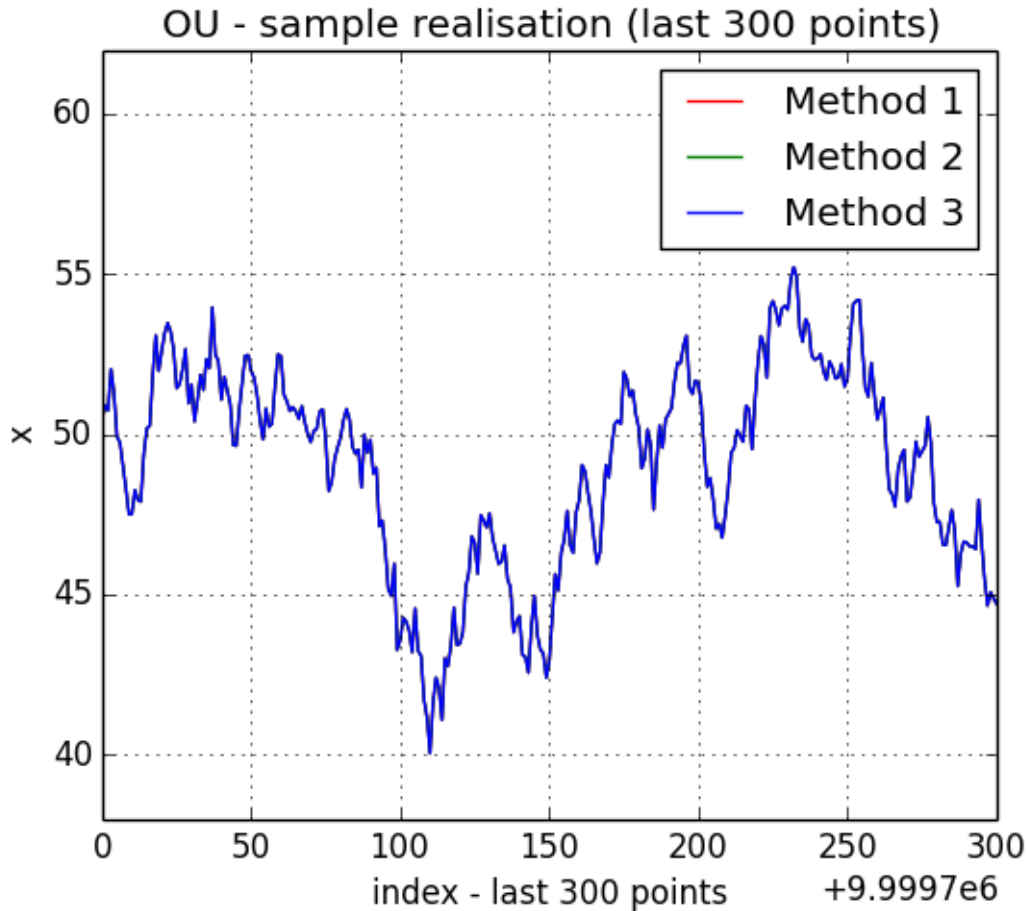
legend(borderaxespad=0.5); # Add a legend to the current figure

```

```

The mean square error comparing x1 and x2 is 0.0
The mean square error comparing x1 and x3 is 1.7638370652718947e-9
The mean square error comparing x2 and x3 is 1.7638370652718947e-9

```



We now analyse (one of) the generated time series, checking for consistency of its statistical properties against the theory. Bear in mind that the larger  $N$ , the more accurate will be the agreement between simulation and theory (i.e. large numbers theorem). We first evaluate mean, variance, and then we estimate the probability density function (i.e. by means of the normalised histogram) and compare it to the theoretical Gaussian curve.

```
In [95]: # Let's use x2, since x1, x2, and x3 are identical
# Note that one should discard a transient because the theoretical PDF is
# However, if x0 is zero the transient should only affect the estimate of
transient      = Int64(floor(10. *  $\tau$  /  $\Delta t$ ) + 1);
x              = x2[transient:end];
N              = length(x)
println("Mean is $(mean(x)) [theoretical  $\mu$ ]");
println("Stdev is $(std(x)) [theoretical  $\sigma$ ]");
```

```
Mean is 49.969602453770236 [theoretical 50.0]
Stdev is 6.014366270070594 [theoretical 6.0]
```

```

In [96]: # Let's now estimate the probability density function (pdf), through the

fig = figure("Probability density function",figsize=(6,5)); # Create a new figure

Npts          = 1000;
(edges, counts) = hist(x, Npts);                          # Compute the histogram of x
xx            = midpoints(edges);                          # Midpoints of the bins with width ΔE

ΔE            = edges[3] - edges[2];                      # Do not use edges[1] as it is not the width
pdf           = counts / (N * ΔE);                        # Normalization, so it becomes a probability

pdf_theory= (1. / sqrt(2 * π * σ^2)) * exp(-(xx-μ).*(xx-μ)/(2 * σ^2)); # Analytical solution

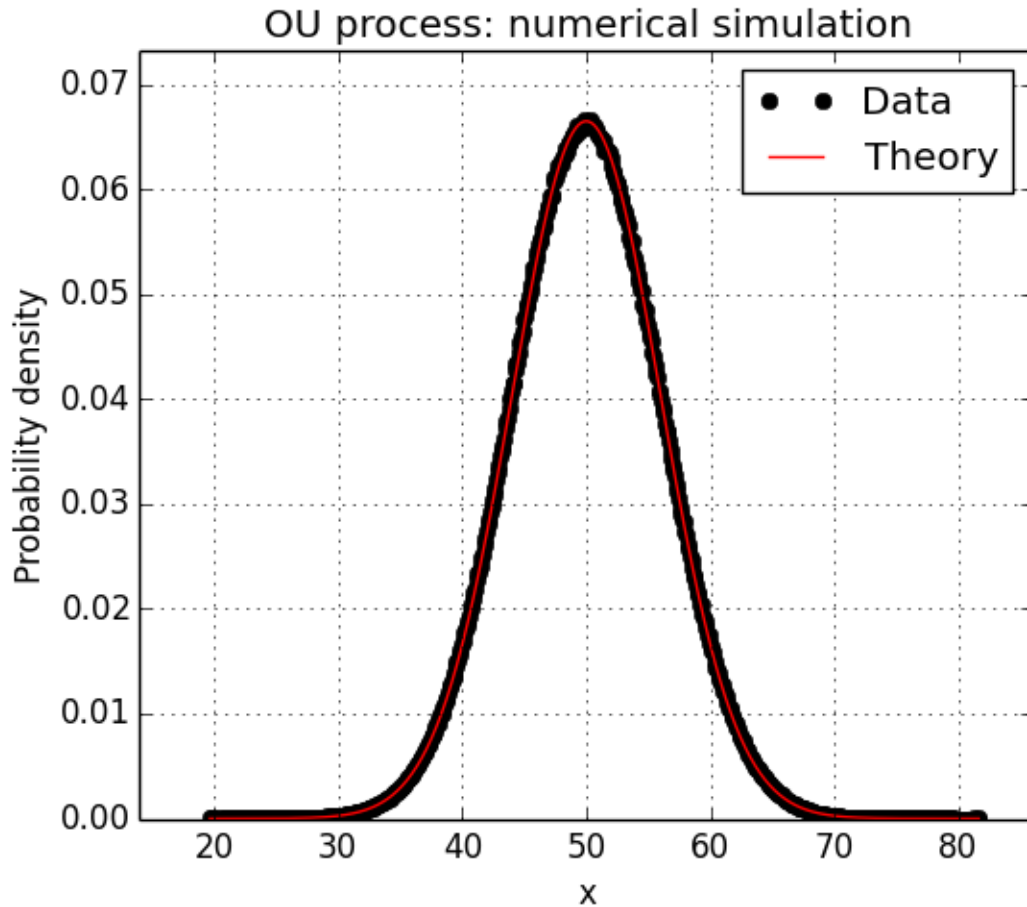
plot(xx, pdf, "ko", label="Data");                        # Plot in the current figure
plot(xx, pdf_theory, "r", label="Theory");                #, LineWidth="2");

ax = gca()                                                 # Get a "handle" on the current axes
ax[:set_xlim]((μ-6*σ), (μ+6*σ));                          # For axes "ax" set the horizontal limits
ax[:set_ylim]((0, 1.1*maximum(pdf)));                     # For axes "ax" set the vertical limits

xlabel("x");                                               # Label of the horizontal axis
ylabel("Probability density");                             # Label of the vertical axis
title("OU process: numerical simulation");                 # Label of the title, at the top
grid("on");                                                # "Grid" on

legend(borderaxespad=0.5);                                # Add a legend to the current figure

```



We finally compute the autocorrelation (i.e. the covariance of the process upon removing the *offset*, being the square mean).

In [97]: *# Let's finally estimate the autocorrelation function (i.e. autocovariance)*

```
fig = figure("OU - autocorrelation", figsize=(6,5));      # Create a new figure

# Let's remove the mean from x, prior to compute the correlation
xc = xcorr(x - mean(x), x - mean(x)); # xcorr is like the corresponding
lags = 1.0 * collect((-N+1):(N-1));    # The vector containing the "lags"

normalised_xc = xc ./ (N-abs(lags));      # Unbiased normalisation
xc_theory      =  $\sigma^2 \exp(-\text{abs}(\text{lags} \cdot \Delta t) / (\tau))$ ; # Theoretical autocorrelation
```

PyPlot.Figure(PyObject <matplotlib.figure.Figure object at 0x32f9aee50>)

```
In [101]: M = 3;                                     # For reducing the points
          plot( $\Delta t \cdot \text{lags}[1:M:\text{end}]$ , normalised_xc[1:M:end], "ko", label="Data"); # Plot
```

```

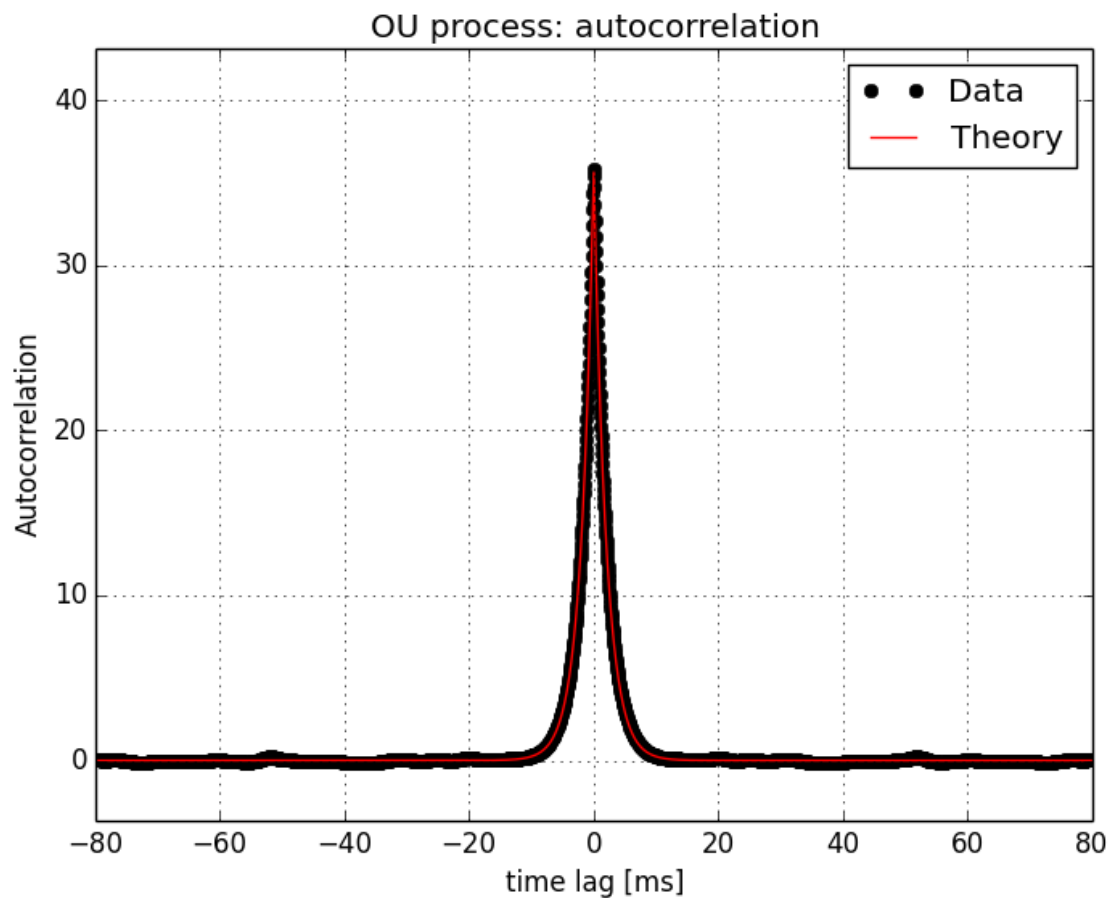
plot( $\Delta t$ *lags[1:M:end], xc_theory[1:M:end], "r", label="Theory");#, LineW

ax = gca()
ax[:set_xlim](-40* $\tau$ , 40* $\tau$ );
ax[:set_ylim]((-0.1* $\sigma^2$ , 1.2* $\sigma^2$ ));

xlabel("time lag [ms]");
ylabel("Autocorrelation");
title("OU process: autocorrelation");
grid("on");

legend(borderaxespad=0.5);

```



In [ ]: